

I build AI products and the systems that make them reliable in production.

8 years building production software inside regulated healthcare — hands-on across services, data, identity, and delivery — and 2.5 years leading the teams doing it. The work spans both halves of a product: the customer-facing platform initiatives and the platform underneath them. Recent independent work is specialised on agentic products and the integration infrastructure they need: persistent agent memory evaluated by ablation, a controlled experiment in agent coordination, and an open standard for repository intelligence. Each ships with the evidence needed to tell whether it works, which is what decides whether an agent product survives production.

Customer-facing healthcare platforms · regulated production

BlueCross BlueShield of Tennessee

2016 – 2026 Chattanooga, TN

Engineering Team Lead

10/2023 – 03/2026

- Aligned high-visibility initiatives including Portal Refresh, CIAM, and Shared Health across security, product, middleware, and engineering teams.
- Optimized authentication workflows with Ping Security and CyberArk, strengthening compliance and system resilience.
- Championed AI-assisted engineering and test automation using Cursor AI, Playwright, and Cypress to reduce manual overhead within corporate AI constraints.
- Led enterprise modernization from legacy Java monoliths to REST-based Node.js microservices, improving developer velocity and platform scalability.

Engineering & DevOps Strategy Lead

08/2022 – 10/2023

- Established backlog governance and prioritized mission-critical work when product ownership capacity was constrained, preventing uncontrolled scope growth.
- Introduced AI-assisted test automation to accelerate release cycles and reduce QA bottlenecks.
- Developed CI/CD frameworks with Azure DevOps and Argo CD for scalable, secure application deployments.

Application Systems Analyst Specialist / Technical Lead

08/2019 – 08/2022

- Architected high-availability microservices to improve reliability and application performance.
- Collaborated with security teams to establish IAM protocols using OAuth and SAML, supporting enterprise compliance.

Application Developer / Analyst · Application Developer

08/2016 – 08/2019

BOUNDARY

Titles, dates, and scope are stated as held. The named initiatives are enterprise programmes this role led or aligned across security, product, middleware, and engineering — not products owned end to end by one engineer. No proprietary metrics are stated.

03 / SELECTED  
AI PRODUCTS

Independent  
product work •  
Marcelle Labs

**Never Ask Twice** Persistent agent memory, evaluated by ablation VERIFY ↗

- Enterprise support MemoryAgent that carries customer context across sessions, with an explicit forgetting policy rather than unbounded accumulation.

**Interlock** Controlled evaluation of agent coordination VERIFY ↗

- Execution separated from observation: frozen evidence packet, independent verifier, explicitly stated non-claims.

**workspace.json** Open standard for repository-derived intelligence VERIFY ↗

- Designed and steward an Apache-2.0 specification — schema, producer tooling, governance — for committing repository intelligence in a deterministic, tool-consumable form.

ALSO **Vreko** — MCP and codebase intelligence for agentic development.

BOUNDARY Public source and recorded evidence establish that these systems exist and what was observed; they do not establish adoption or production scale outside the repositories where they have been applied.

04 / AGENT  
PLATFORM

Linear as an  
execution control  
plane

**META-268 Native delegation versus a custom orchestrator** Decided • boundary held

Native delegation was kept; the custom boundary was drawn around what Linear does not own — local execution, repository context, and the evidence written back.

**META-331 Linear as the control plane for execution surfaces** Implemented • in use

An execution-surface responsibility model with Linear as the control plane: the issue is the one place work is assigned, tracked, and closed, and runners stay interchangeable beneath it.

**INFRA-11 Structured agent lifecycle: thought → progress → completion** Implemented

Agent lifecycle integration structured as thought → progress → completion, so a session leaves a readable trace on the issue rather than one terminal comment.

BOUNDARY Curated summaries of private workspace decisions, written for publication rather than exported. No public artifact stands behind them: stated claims, not verified evidence.

05 / TECHNICAL  
FOUNDATION

CAPABILITIES

|                     |                                                          |
|---------------------|----------------------------------------------------------|
| PRODUCT / WEB       | TypeScript · Node.js · REST APIs · Playwright · Cypress  |
| AI / AGENTS         | MCP · JSON-RPC · tool contracts · Google ADK · Vertex AI |
| BACKEND / DATA      | Java · microservices · JSON Schema · Postgres / pgvector |
| PLATFORM / DELIVERY | Azure DevOps · Tekton · Argo CD · OpenShift · Cloud Run  |
| IDENTITY / SECURITY | OAuth · SAML · IAM (Ping, CyberArk) · CIAM               |

|           |                                                   |                               |             |
|-----------|---------------------------------------------------|-------------------------------|-------------|
| EDUCATION | <b>Master's Degree, Organizational Leadership</b> | Saint Joseph's University     | 2019 – 2024 |
|           | <b>Bachelor of Science, Computer Science</b>      | Southern Adventist University | 2009 – 2015 |